

```
In [1]: #Importing pandas library to get some information about dataset
```

```
import pandas as pd
```

```
# Loading the training and testing data of Titanic datasets
```

```
train_data = pd.read_csv('train.csv')
```

```
test_data = pd.read_csv('test.csv')
```

```
In [2]: # exploring the dataset
```

```
print(train_data.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
In [3]: print(train_data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 891 entries, 0 to 890
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	object
4	Sex	891 non-null	object
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64
7	Parch	891 non-null	int64
8	Ticket	891 non-null	object
9	Fare	891 non-null	float64
10	Cabin	204 non-null	object
11	Embarked	889 non-null	object

```
dtypes: float64(2), int64(5), object(5)
```

```
memory usage: 83.7+ KB
```

```
None
```

```
In [4]: # Dropping unnecessary columns or handling missing values in the dataset
```

```
# here, dropping 'Name', 'Ticket', 'Cabin' columns which might not directly impact
```

```
train_data.drop(['Name', 'Ticket', 'Cabin'], axis=1, inplace=True)
```

```
test_data.drop(['Name', 'Ticket', 'Cabin'], axis=1, inplace=True)
```

```
In [5]: #Handling missing values (e.g., filling missing age values with median)
```

```
train_data['Age'].fillna(train_data['Age'].median(), inplace=True)
test_data['Age'].fillna(test_data['Age'].median(), inplace=True)
test_data['Fare'].fillna(test_data['Fare'].median(), inplace=True)
```

```
In [6]: # Converting categorical variables into numerical ones (e.g., 'Sex', 'Embarked')
```

```
train_data = pd.get_dummies(train_data, columns=['Sex', 'Embarked'])
test_data = pd.get_dummies(test_data, columns=['Sex', 'Embarked'])
```

```
In [7]: # Defining features and target variable
```

```
features = train_data.drop('Survived', axis=1)
target = train_data['Survived']
```

```
In [8]: from sklearn.model_selection import KFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```
# Initializing the model with random forest classifier
```

```
model = RandomForestClassifier()
```

```
# Define the number of folds we want
```

```
num_folds = 5
```

```
# Initializing KFold
```

```
kfold = KFold(n_splits=num_folds, shuffle=True, random_state=42)
```

```
In [15]: # Perform cross-validation
```

```
fold_accuracy = []
```

```
for fold, (train_idx, val_idx) in enumerate(kfold.split(features, target)):
    X_train, X_val = features.iloc[train_idx], features.iloc[val_idx]
    y_train, y_val = target.iloc[train_idx], target.iloc[val_idx]
```

```
# Train the model
```

```
model.fit(X_train, y_train)
```

```
# Make predictions on validation set
```

```
predictions = model.predict(X_val)
```

```
# Calculate accuracy for this fold
```

```
accuracy = accuracy_score(y_val, predictions)
```

```
fold_accuracy.append(accuracy)
```

```
print(f"Fold {fold+1} Accuracy: {accuracy}")
```

```
Fold 1 Accuracy: 0.8268156424581006
```

```
Fold 2 Accuracy: 0.7921348314606742
```

```
Fold 3 Accuracy: 0.848314606741573
```

```
Fold 4 Accuracy: 0.7808988764044944
```

```
Fold 5 Accuracy: 0.8258426966292135
```

```
In [ ]:
```

```
In [16]: # Calculate average accuracy across all folds
```

```
average_accuracy = sum(fold_accuracy) / len(fold_accuracy)
```

```
print(f"Average Accuracy: {average_accuracy}")
```

```
Average Accuracy: 0.8148013307388112
```

```
In [17]: from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import classification_report
```

```

# Splitting the training data into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(features, target, test_size=0.2,

# Initialize a simple model (e.g., RandomForestClassifier)
simple_model = RandomForestClassifier()

# Train the model using the training set
simple_model.fit(X_train, y_train)

# Validate the model using the validation set
val_predictions = simple_model.predict(X_val)

# Evaluate the model's performance on the validation set
val_accuracy = accuracy_score(y_val, val_predictions)
print(f"Validation Accuracy: {val_accuracy}")

# Generate a classification report for detailed performance analysis
print("Classification Report:")
print(classification_report(y_val, val_predictions))

```

Validation Accuracy: 0.8212290502793296

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.89	0.85	105
1	0.82	0.73	0.77	74
accuracy			0.82	179
macro avg	0.82	0.81	0.81	179
weighted avg	0.82	0.82	0.82	179

```

In [18]: # Using RandomForestClassifier to get feature importances
feature_importances = model.feature_importances_

# Identify and select important features
important_features = pd.Series(feature_importances, index=features.columns).sort_values(ascending=False)
selected_features = important_features[:8].index # Select top 8 important features

# Retrain the model using selected features
selected_features_model = RandomForestClassifier()
selected_features_model.fit(X_train[selected_features], y_train)

# Validate the model using the selected features
val_predictions_selected = selected_features_model.predict(X_val[selected_features])

# Evaluate the model's performance on the validation set
val_accuracy_selected = accuracy_score(y_val, val_predictions_selected)
print(f"Validation Accuracy with Selected Features: {val_accuracy_selected}")

```

Validation Accuracy with Selected Features: 0.8156424581005587

```

In [19]: from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline

# Apply regularization with a Logistic Regression model
scaler = StandardScaler()
regularized_model = make_pipeline(scaler, LogisticRegression(penalty='l1', solver='libsvm'))
regularized_model.fit(X_train, y_train)

# Validate the model using regularization
val_predictions_regularized = regularized_model.predict(X_val)

```

```
# Evaluate the model's performance on the validation set
val_accuracy_regularized = accuracy_score(y_val, val_predictions_regularized)
print(f"Validation Accuracy with Regularization: {val_accuracy_regularized}")
```

Validation Accuracy with Regularization: 0.8044692737430168

```
In [20]: from sklearn.ensemble import GradientBoostingClassifier

# Experimenting with GradientBoostingClassifier
gradient_boost_model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1)
gradient_boost_model.fit(X_train, y_train)

# Validating the model using Gradient Boosting Classifier
val_predictions_gradient_boost = gradient_boost_model.predict(X_val)

# Evaluating the model's performance on the validation set
val_accuracy_gradient_boost = accuracy_score(y_val, val_predictions_gradient_boost)
print(f"Validation Accuracy with GradientBoostingClassifier: {val_accuracy_gradient_boost}")
```

Validation Accuracy with GradientBoostingClassifier: 0.8156424581005587

```
In [21]: from sklearn.svm import SVC

# Experimenting with Support Vector Machines
svm_model = SVC(kernel='linear', C=1.0, random_state=42)
svm_model.fit(X_train, y_train)

# Validate the model using Support Vector Machine
val_predictions_svm = svm_model.predict(X_val)

# Evaluate the model's performance on the validation data set
val_accuracy_svm = accuracy_score(y_val, val_predictions_svm)
print(f"Validation Accuracy with SVM: {val_accuracy_svm}")
```

Validation Accuracy with SVM: 0.776536312849162

```
In [9]: ###Experiments
```

```
In [10]: #####1 Experiment 1: Feature Engineering and Model Comparison
```

```
In [22]: # Creating a 'FamilySize' feature by combining SibSp and Parch
train_data['FamilySize'] = train_data['SibSp'] + train_data['Parch']
test_data['FamilySize'] = test_data['SibSp'] + test_data['Parch']

# Drop SibSp and Parch as they are now redundant
train_data.drop(['SibSp', 'Parch'], axis=1, inplace=True)
test_data.drop(['SibSp', 'Parch'], axis=1, inplace=True)

# Define features and target with new feature
features_with_family = train_data.drop('Survived', axis=1)
target = train_data['Survived']

# Performing train-test split to split data
X_train_with_family, X_val_with_family, y_train, y_val = train_test_split(features_with_family, target, test_size=0.2, random_state=42)
```

```
In [12]: from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression

# Train and evaluate different models using new features
models = {
    'RandomForest': RandomForestClassifier(random_state=42),
    'LogisticRegression': LogisticRegression(random_state=42)
}
```

```

for model_name, model in models.items():
    model.fit(X_train_with_family, y_train)
    predictions = model.predict(X_val_with_family)
    accuracy = accuracy_score(y_val, predictions)
    print(f"{model_name} Accuracy: {accuracy}")

```

RandomForest Accuracy: 0.8324022346368715

LogisticRegression Accuracy: 0.7988826815642458

C:\Users\User\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

In [23]: *###Experiment 2 - Hyperparameter Tuning and Ensemble Methods*

In [24]: **from** sklearn.model_selection **import** GridSearchCV

Hyperparameter tuning for RandomForestClassifier

```

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10]
}

```

```

rf_model = RandomForestClassifier(random_state=42)
grid_search = GridSearchCV(rf_model, param_grid, cv=5)
grid_search.fit(X_train_with_family, y_train)

```

```

best_rf_model = grid_search.best_estimator_
best_rf_model.fit(X_train_with_family, y_train)
rf_predictions = best_rf_model.predict(X_val_with_family)
rf_accuracy = accuracy_score(y_val, rf_predictions)
print(f"RandomForest Tuned Accuracy: {rf_accuracy}")

```

RandomForest Tuned Accuracy: 0.8044692737430168

In [15]: **from** sklearn.ensemble **import** VotingClassifier

Ensemble using VotingClassifier with the best-performing models

```

voting_model = VotingClassifier(
    estimators=[('RandomForest', best_rf_model), ('LogisticRegression', models['Log
    voting='hard'
)

```

```

voting_model.fit(X_train_with_family, y_train)
voting_predictions = voting_model.predict(X_val_with_family)
voting_accuracy = accuracy_score(y_val, voting_predictions)
print(f"Voting Classifier Accuracy: {voting_accuracy}")

```

Voting Classifier Accuracy: 0.8044692737430168

C:\Users\User\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

